

Operationalizing an AWS Machine Learning Project

- 1) Write a short justification for why you chose the SageMaker notebook instance type you did (consider cost computing power, and speed of launching for each instance type)

Answer: I used the **ml.t3.medium** instance which is a general purpose instance. It provides average computing power and speed at low cost. Considering the nature of this project, there was no need for extra high performance computing, hence the choice of this instance to save cost.

- 2) After launching your SageMaker instance, take a screenshot of your sagemaker dashboard's Notebooks > Instances section to show what you have done.

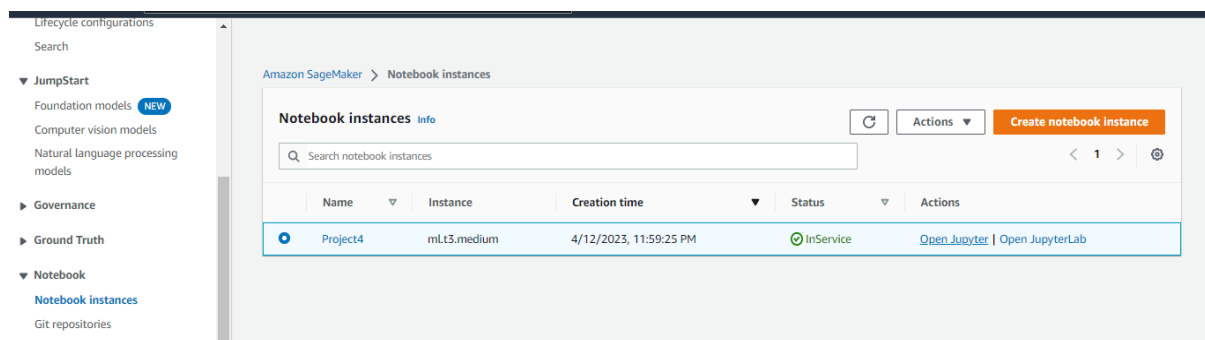


Fig 1: SageMaker notebook instance

- 3) Take a screenshot of your deployed endpoint (single instance training) and note its name

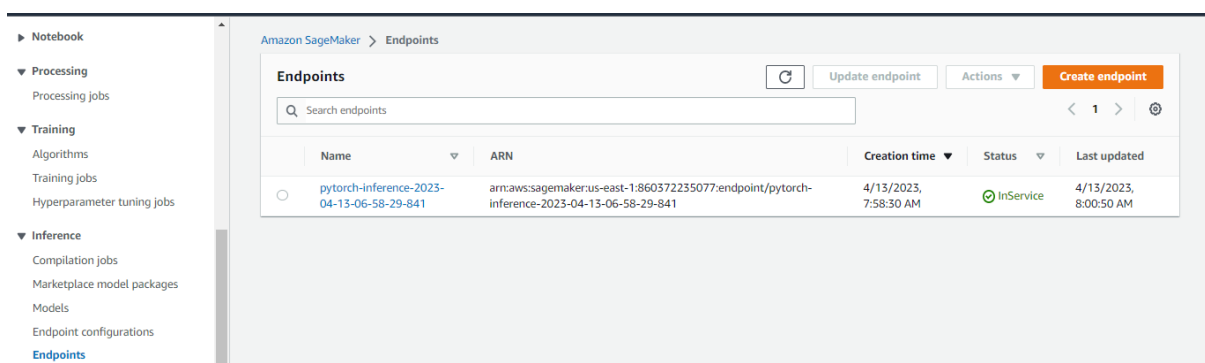


Fig 2: Deployed endpoint after single instance training

- 4) Write a justification for the type of EC2 instance you created.

Answer: Initially, I used the t2.medium instance to save cost. However, when I tried to run the solution.py file, I encountered a numpy import error and I tried to address it by activating pytorch in EC2. The EC2 t2.medium instance was unable to activate pytorch and I switched to a g3s.xlarge instance as suggested by the error message. Relevant screenshots are attached below.

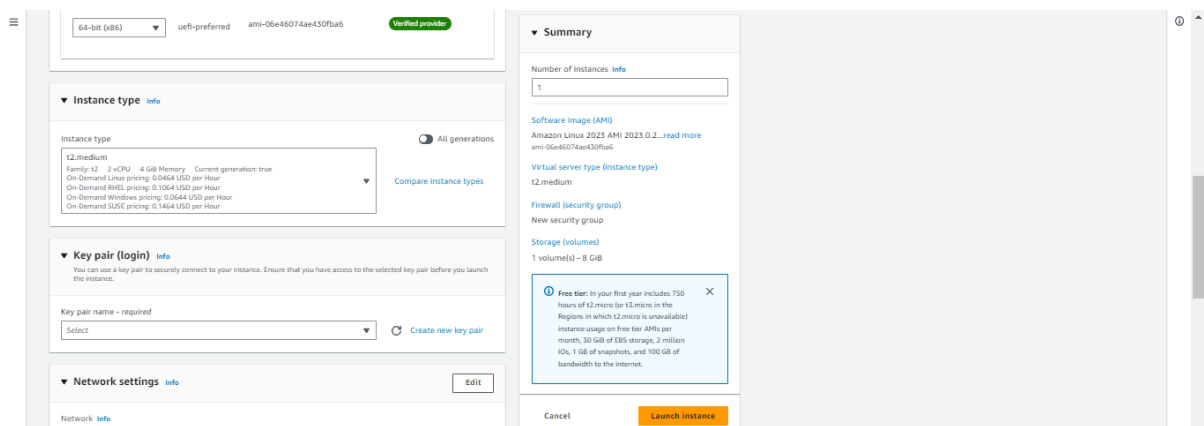


Fig 3: Initial EC2 instance

```
[root@ip-172-31-17-183 ~]# vim solution.py
[root@ip-172-31-17-183 ~]# python solution.py
[root@ip-172-31-17-183 ~]# vim solution.py
[root@ip-172-31-17-183 ~]# python solution.py
Traceback (most recent call last):
  File "solution1.py", line 1, in <module>
    import numpy as np
ImportError: Module named numpy
[root@ip-172-31-17-183 ~]# source activate pytorch
ERROR: Please note that the Amazon EC2 t2.medium instance type is not supported by current Deep Learning AMI.
Please try one of the supported EC2 instances: G3, P3, P3dn, P4d, P4de, G5, G4dn.
Please refer the DLI release notes https://aws.amazon.com/deep-learning-ami-gpu-pytorch-2-0-0-aws-ubuntu-20-04-ami/ for more information.
(pytorch) [root@ip-172-31-17-183 ~]#
```

Fig 4: EC2 error message

- 5) Open the TrainedModels directory in your EC2 instance and take a screenshot of the model that has been saved in it.

```
warnings.warn(msg)
Starting Model Training
base
(pytorch) ubuntu@ip-172-31-34-183:~$ ls TrainedModels
TrainedModels: command not found
(pytorch) ubuntu@ip-172-31-34-183:~$ ls TrainedModels
model.pth
(pytorch) ubuntu@ip-172-31-34-183:~$
```

Fig 5: Model in the TrainedModels directory

- 6) Write on the differences between the code in “train_and_deploy-solution.ipynb” and “ec2train1.py”

Answer: In `ec2train1.py`, the parameters, estimators, and tuners are not explicitly defined. `ec2train1.py` is optimized by building functions to handle model training and testing. These functions can be called on data contained in EC2 when needed. In the `train_and_deploy` solution notebook, however, the data is imported and model training and testing is done using explicitly defined parameters, estimators, and tuners. In addition, libraries such as `boto3` and `sagemaker` cannot be used in `ec2train1.py`.

- 7) Take a screenshot of the lambda function after running the test event, and provide the list of numbers given in the output.



Fig 6: Lambda function test output

```
[
  [-0.3194071054458618, 0.06982807070016861, -0.10188700258731842,
    0.05443248152732849, 0.23411345481872559, 0.2769347131252289,
    0.016361981630325317, 0.09543728828430176, -0.23329131305217743,
    0.1326563060283661, 0.08965423703193665, 0.0579850859940052,
    0.15090619027614594, 0.10275612771511078, 0.1422945261001587,
    0.02920246124267578, -0.02602965384721756, -0.39369359612464905,
    0.20661674439907074, 0.05201426148414612, 0.020475119352340698,
    0.1092258021235466, 0.14406603574752808, 0.16800859570503235,
    0.2569163739681244, -0.4271831512451172, 0.18519136309623718,
    0.3146352469921112, 0.034361548721790314, -0.2248910516500473,
    0.07397255301475525, 0.01564362645149231, -0.2630861699581146,
    0.2257373183965683, -0.19176961481571198, -0.13752171397209167,
    0.032651618123054504, -0.15763114392757416, -0.059963926672935486,
    0.03800392150878906, -0.2107524275779724, -0.0781690998655319,
    0.10200075805187225, -0.35978806018829346, 0.0466545969247818,
    0.13159333169460297, 0.02251681685447693, 0.10969046503305435,
    0.12295417487621307, -0.048323310911655426, -0.026371128857135773,
    0.07860969007015228, -0.2702891230583191, 0.07395516335964203,
    0.31811490654945374, 0.15498119592666626, 0.020910531282424927,
    0.13895606994628906, -0.1581280380487442, 0.056791454553604126,
    0.16812896728515625, -0.0019059106707572937, -0.36973121762275696,
    0.4153045117855072, -0.12665320932865143, -0.47466224431991577,
    0.034660764038562775, -0.3707352876663208, -0.09482190012931824,
    0.1700029969215393, 0.22720716893672943, -0.180681973695755,
    0.30132168531417847, -0.1123449057340622, -0.23300659656524658,
    0.09635794162750244, -0.29833337664604187, -0.25851720571517944,
    0.12416891753673553, -0.09933248162269592, -0.031318724155426025,
    0.015487030148506165, -0.10894891619682312, 0.1430465281009674,
  ]
]
```

```

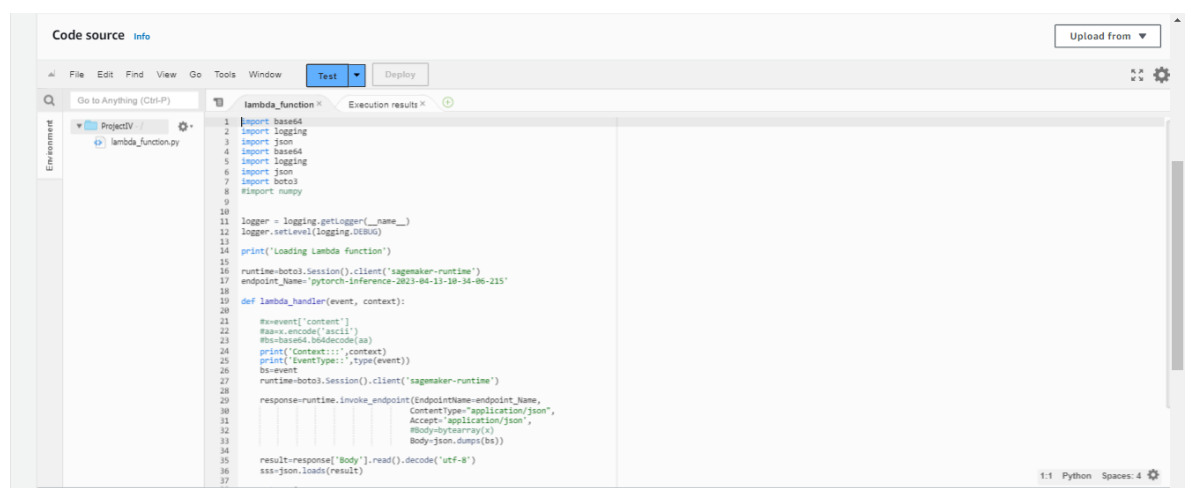
0.24369531869888306, -0.22964990139007568, -0.17065121233463287, -
0.1537196934223175, -0.10664661228656769, -0.0409412682056427, -
0.030008748173713684, 0.006851717829704285, -0.4325335621833801, -
0.12947550415992737, -0.056942954659461975, -0.1904909610748291, -
0.17198391258716583, 0.133949875831604, -0.1790093183517456, -
0.1203477531671524, 0.06557272374629974, -0.41081827878952026, -
0.039602264761924744, -0.3460289239883423, -0.27233603596687317, -
0.06590510904788971, -0.17482388019561768, -0.1722773313522339, -
0.02484789490699768, -0.3098023235797882, -0.17887285351753235, -
0.005011945962905884, -0.15574386715888977, -0.27694201469421387, -
0.11342388391494751, -0.2868531346321106, 0.013332098722457886, -
0.19056998193264008, -0.059050947427749634, -0.0750817060470581, -
0.37999147176742554, -0.29933661222457886, -0.0784531682729721, -
0.05784786492586136, -0.36373087763786316, -0.18035843968391418, -
0.1644153892993927, -0.4166147708892822, -0.03248770534992218, -
0.3103395402431488, -0.4522474408149719, -0.25467294454574585, -
0.284565806388855]

```

Fig 7: List of numbers in Lambda function test output

8) Describe how the Lambda function is written and how it works

Answer: The `lambda_handler` function is the major component of this Lambda function. It takes in a picture (passed in as a url using the JSON file of the Lambda function) which it passes through the model contained in a SageMaker endpoint (the endpoint is defined within the Lambda function as `endpoint_Name`). The model contained in this specified endpoint is a dog breed classifier. Therefore, the output of this Lambda function is a list containing the probability of the dog in the input picture belonging to the various breeds in the dataset.



```

1 import boto3
2 import logging
3 import json
4 import boto3
5 import logging
6 import json
7 import boto3
8 import numpy
9
10
11 logger = logging.getLogger(__name__)
12 logger.setLevel(logging.DEBUG)
13
14 print('Loading Lambda function')
15
16 runtime=boto3.Session().client('sagemaker-runtime')
17 endpoint_Name='pytorch-inference-2023-04-13-10-34-06-215'
18
19 def lambda_handler(event, context):
20
21     #event['content']
22     #aws.encode('ascii')
23     #bs=boto3.Session().client('sagemaker-runtime')
24     print('Context:', context)
25     print('Event type:', type(event))
26     bs=boto3.Session().client('sagemaker-runtime')
27
28     response=runtime.invoke_endpoint(EndpointName=endpoint_Name,
29                                     ContentType='application/json',
30                                     Accept='application/json',
31                                     Body=json.dumps(event))
32
33     result=response['Body'].read().decode('utf-8')
34
35     sss=json.loads(result)
36
37

```

Fig 8: Lambda function

9) Take a screenshot of the policies attached to your Lambda function role

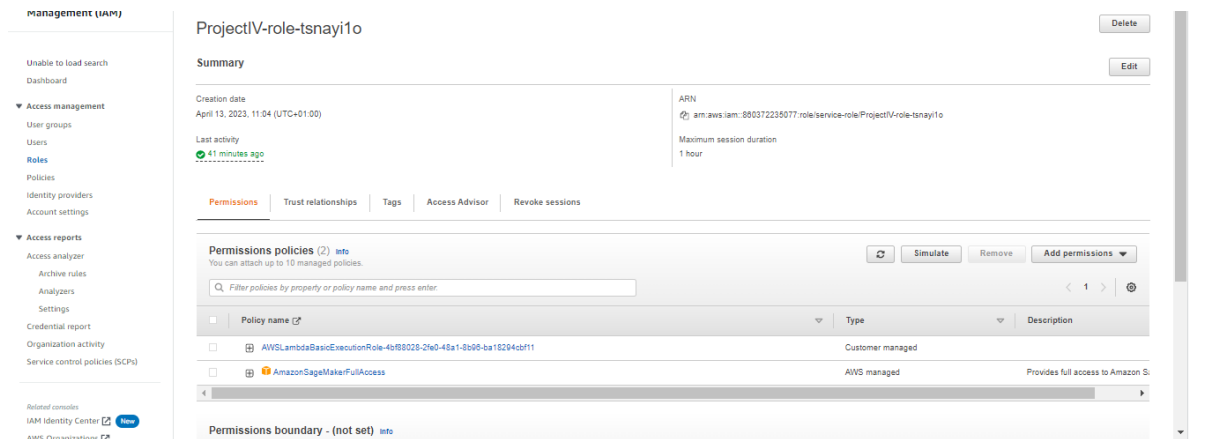


Fig 9: Policies attached to the Lambda function role before security considerations

10) Write about the security vulnerabilities in your AWS workspace

Answer: Attaching the AmazonSageMakerFullAccess policy to my Lambda function exposes my account to security risks through the function. This is because the Lambda function has access to all my endpoints in SageMaker.

To address this issue, I replaced the AmazonSageMakerFullAccess policy with a policy that allows my Lambda function access to only the pytorch inference endpoint that it passes data to. Appropriate screenshots are attached below.

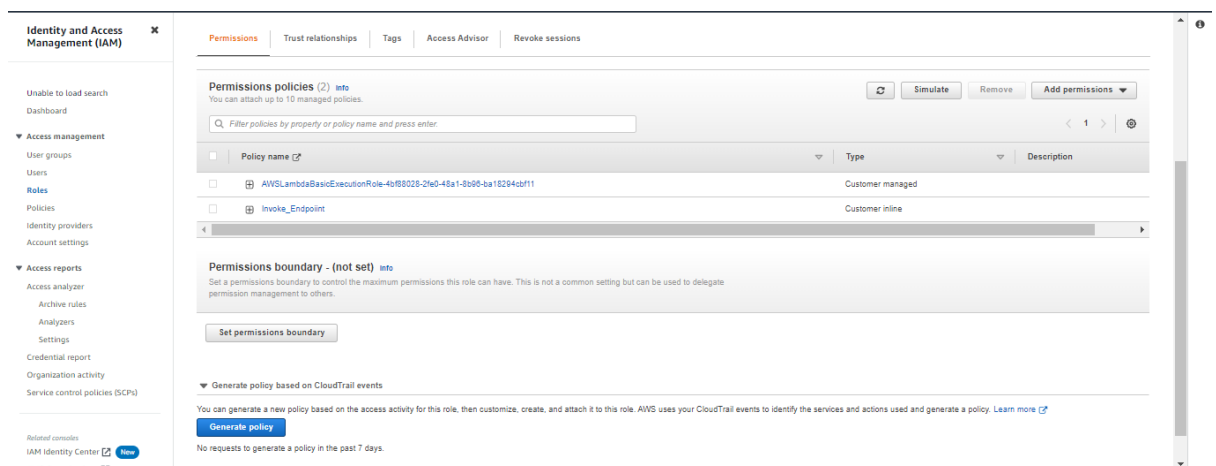


Fig 10: Policies attached to the Lambda function role after security considerations

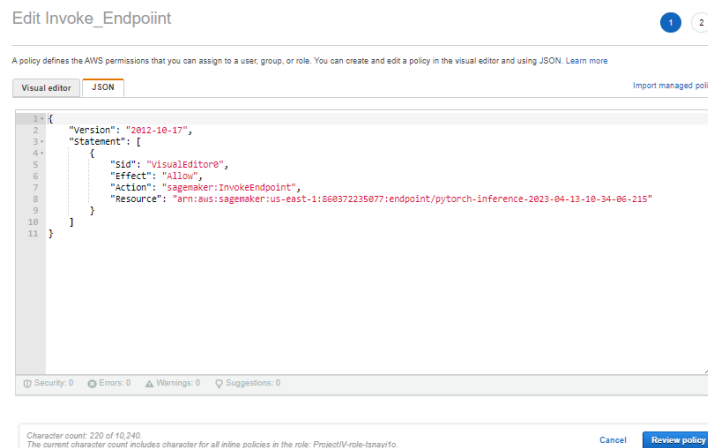


Fig 11: JSON for new policy attached to Lambda function

11) Write about the choices you made during concurrency and autoscaling setup

Answer:

Autoscaling: I set the target value to 40. This means that instances will be added when there are greater than or equal to 40 simultaneous requests to this endpoint. I set the value high to ensure that new instances are not activated unless absolutely necessary, ultimately leading to less cost.

I set the maximum instance count to 2 because I do not expect high traffic on my endpoint.

I set 30 for the scale in cool down value to ensure that extra resources are not employed for momentary spikes in request traffic. 30 minutes is enough time to ascertain that there is indeed high traffic, before deploying more resources.

I also set 30 for the scale out cool down value to ensure that resources are not retrieved for temporary reductions in request traffic. 30 minutes is sufficient time to ensure that there is indeed traffic reduction before retrieving deployed resources.

Concurrency: Since I do not expect high traffic on this project and provisioned concurrency needs to be lower than reserved concurrency, I picked 5 reserved instances and 3 provisioned instances to save cost.

Foundation models **NEW**

Computer vision models

Natural language processing models

► Governance

► Ground Truth

▼ Notebook

Notebook instances

Git repositories

► Processing

▼ Training

Algorithms

Training jobs

Hyperparameter tuning jobs

▼ Inference

Compilation jobs

Marketplace model packages

Models

Endpoint configurations

Endpoints

Batch transform jobs

Shadow tests

IAM role

Amazon SageMaker uses the following service-linked role for automatic scaling. [Learn more](#)

`AWSServiceRoleForApplicationAutoScaling_SageMakerEndpoint`

Built-in scaling policy [Learn more](#)

Policy name

SageMakerEndpointInvocationScalingPolicy

Target metric

SageMakerVariantInvocationsPerInstance [Learn more](#)

Target value

40

Scale in cool down (seconds) - optional

30

Scale out cool down (seconds) - optional

30

☐ Disable scale in

Select if you don't want automatic scaling to delete instances when traffic decreases. [Learn more](#)

Custom scaling policy [Learn more](#)

There are no custom scaling policies for this variant.

Cancel

Save

Fig 12: Autoscaling configuration